



# MAPREDUCE FUNDAMENTAL CONCEPTS



# Why MapReduce?

- Distributes the processing of data on your cluster
- Divides your data up into partitions that are MAPPED (transformed) and REDUCED (aggregated) by mapper and reducer functions you define
- Resilient to failure - an application master monitors your mappers and reducers on each partition



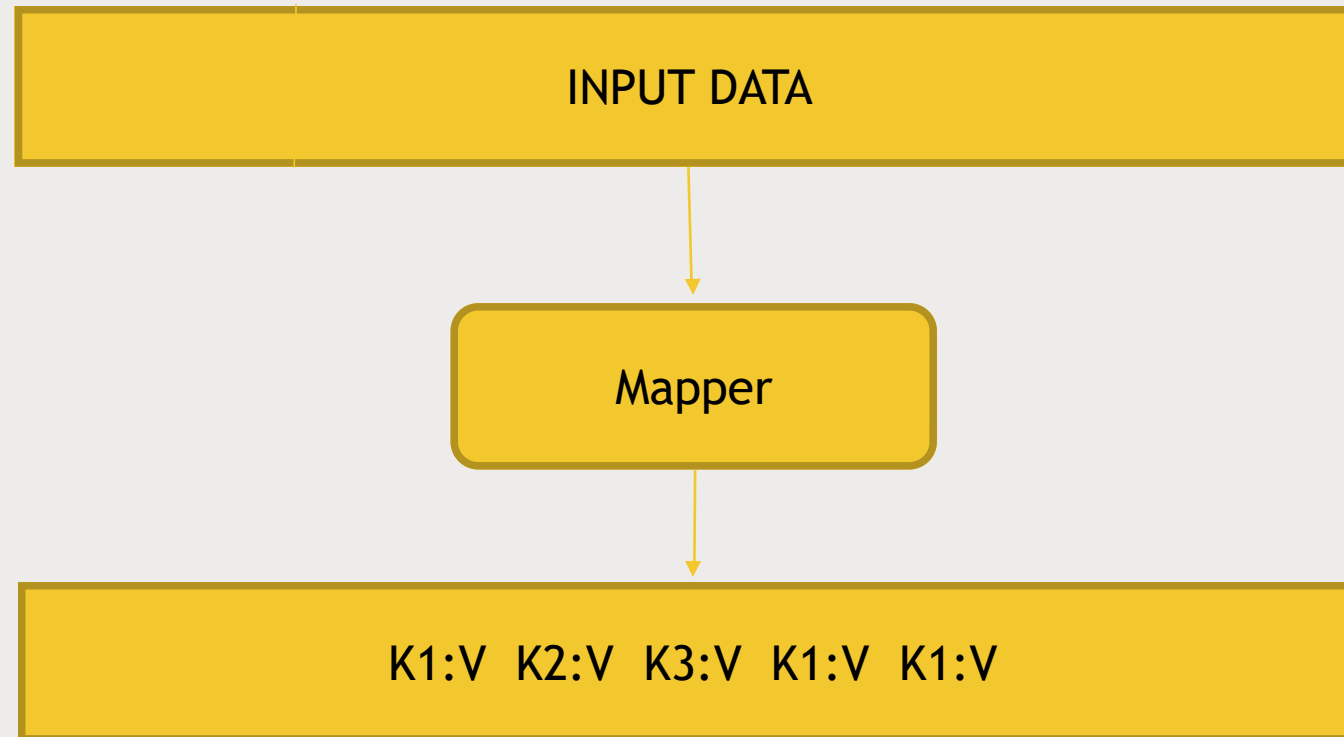
# Let's illustrate with an example

- How many movies did each user rate in the MovieLens data set?



# How MapReduce Works: Mapping

- The MAPPER converts raw source data into **key/value** pairs



# Example: MovieLens Data (u.data file)

USER ID | MOVIE ID | RATING | TIMESTAMP

196 242 3 881250949

186 302 3 891717742

196 377 1 878887116

244 51 2 880606923

166 346 1 886397596

186 474 4 884182806

186 265 2 881171488

# Map users to movies they watched

| USER ID | MOVIE ID | RATING | TIMESTAMP |
|---------|----------|--------|-----------|
|---------|----------|--------|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 242 | 3 | 881250949 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 302 | 3 | 891717742 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 377 | 1 | 878887116 |
|-----|-----|---|-----------|

|     |    |   |           |
|-----|----|---|-----------|
| 244 | 51 | 2 | 880606923 |
|-----|----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 166 | 346 | 1 | 886397596 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 474 | 4 | 884182806 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 265 | 2 | 881171488 |
|-----|-----|---|-----------|

Mapper

196:242 186:302 196:377 244:51 166:346 186:274 186:265

# Extract and Organize What We Care About

196:242 186:302 196:377 244:51 166:346 186:274 186:265



# MapReduce Sorts and Groups the Mapped Data (“Shuffle and Sort”)

196:242 186:302 196:377 244:51 166:346 186:274 186:265



166:346    186:302,274,265    196:242,377    244:51



# The REDUCER Processes Each Key's Values

166:346    186:302,274,265    196:242,377    244:51



len(movies)



166:1    186:3    196:2    244:1

# Putting it All Together

| USER ID | MOVIE ID | RATING | TIMESTAMP |
|---------|----------|--------|-----------|
|---------|----------|--------|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 242 | 3 | 881250949 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 302 | 3 | 891717742 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 377 | 1 | 878887116 |
|-----|-----|---|-----------|

|     |    |   |           |
|-----|----|---|-----------|
| 244 | 51 | 2 | 880606923 |
|-----|----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 166 | 346 | 1 | 886397596 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 474 | 4 | 884182806 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 265 | 2 | 881171488 |
|-----|-----|---|-----------|



**MAPPER**



196:242 186:302 196:377 244:51 166:346 186:274 186:265



**SHUFFLE AND SORT**



166:346 186:302,274,265 196:242,377 244:51



**REDUCER**



166:1 186:3 196:2 244:1

A decorative L-shaped frame in a dark blue color, consisting of a horizontal bar at the top left and a vertical bar on the right side, both extending towards the bottom right corner.

# MAPREDUCE ON A CLUSTER

How MapReduce Scales

# Putting it All Together

| USER ID | MID | R | T |
|---------|-----|---|---|
|---------|-----|---|---|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 242 | 3 | 881250949 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 302 | 3 | 891717742 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 196 | 377 | 1 | 878887116 |
|-----|-----|---|-----------|

|     |    |   |           |
|-----|----|---|-----------|
| 244 | 51 | 2 | 880606923 |
|-----|----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 166 | 346 | 1 | 886397596 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 474 | 4 | 884182806 |
|-----|-----|---|-----------|

|     |     |   |           |
|-----|-----|---|-----------|
| 186 | 265 | 2 | 881171488 |
|-----|-----|---|-----------|

MAPPER

196:242 186:302

196:377 244:51

166:346 186:274 186:265

SHUFFLE AND SORT

166:346 186:302,274,265

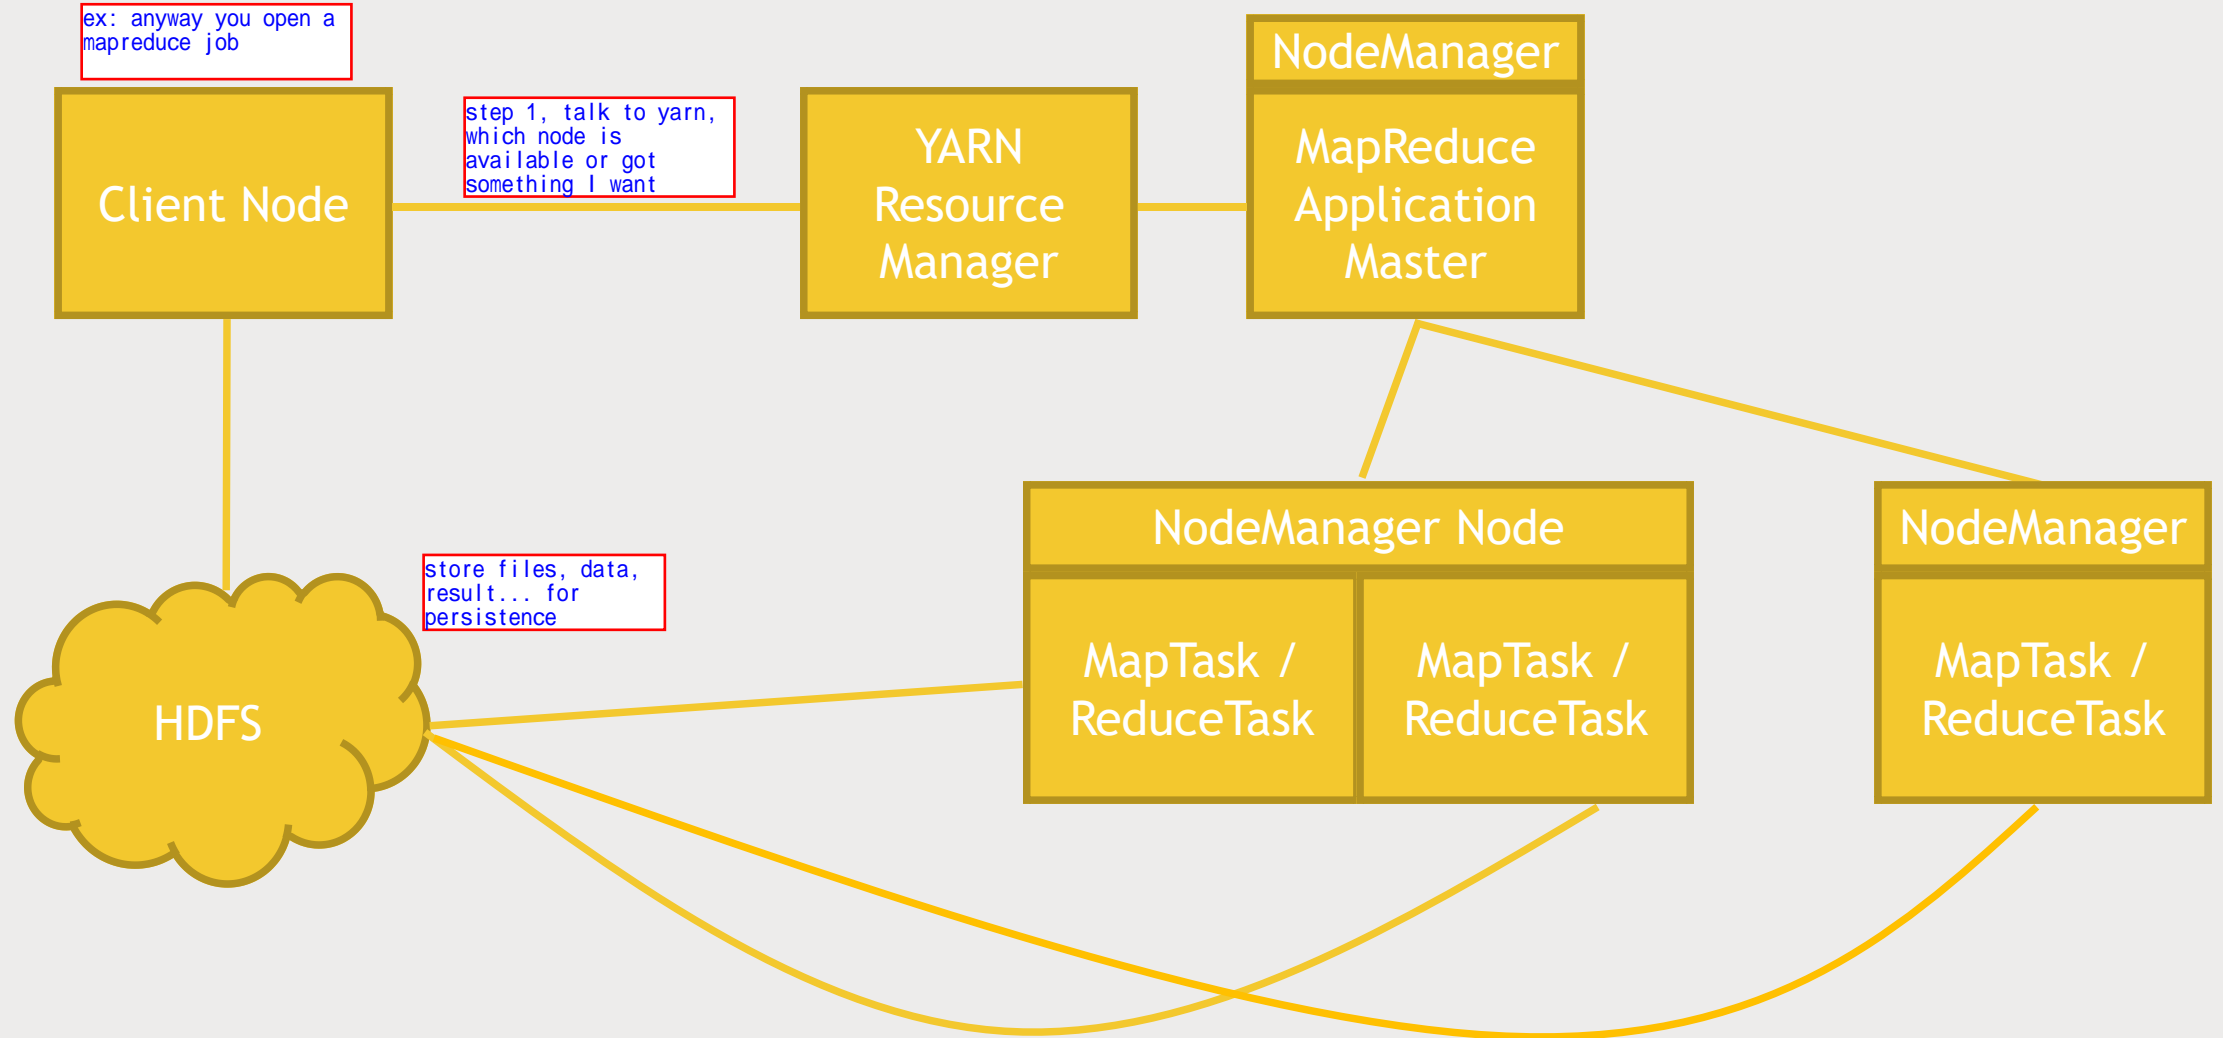
196:242,377 244:51

REDUCER

166:1 186:3

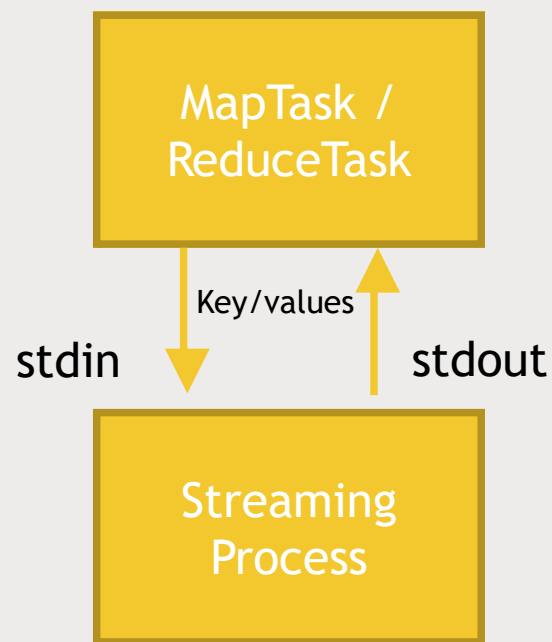
196:2 244:1

# What's Happening



# How are mappers and reducers written?

- MapReduce is natively Java
- STREAMING allows interfacing to other languages (ie Python)




# Handling Failure



- Application master monitors worker tasks for errors or hanging
  - *Restarts as needed*
  - *Preferably on a different node*
- What if the application master goes down?
  - *YARN can try to restart it*
- What if an entire Node goes down?
  - *This could be the application master*
  - *The resource manager will try to restart it*
- What if the resource manager goes down?
  - *Can set up “high availability” (HA) using Zookeeper to have a hot standby*





# MAPREDUCE: A REAL EXAMPLE

How many of each rating type exist?





# How many of each movie rating exist?



# Making it a MapReduce problem

- MAP each input line to (rating, 1)
- REDUCE each rating with the sum of all the 1's

USER ID|MOVIE ID|RATING|TIMESTAMP

196 242 3 881250949

186 302 3 891717742

196 377 1 878887116

244 51 2 880606923

166 346 1 886397596

186 474 4 884182806

186 265 2 881171488

Map

3,1

3,1

1,1

2,1

1,1

4,1

2,1

Shuffle  
& Sort

1 -> 1, 1

2 -> 1, 1

3 -> 1, 1

4 -> 1

Reduce

1, 2

2, 2

3, 2

4, 1

# Writing the Mapper

USER ID|MOVIE ID|RATING|TIMESTAMP

196 242 3 881250949  
186 302 3 891717742  
196 377 1 878887116  
244 51 2 880606923  
166 346 1 886397596  
186 474 4 884182806  
186 265 2 881171488

Map

3,1  
3,1  
1,1  
2,1  
1,1  
4,1  
2,1

Shuffle  
& Sort

1 -> 1, 1  
2 -> 1, 1  
3 -> 1, 1  
4 -> 1

Reduce

1, 2  
2, 2  
3, 2  
4, 1

```
def mapper_get_ratings(self, _, line):  
    (userID, movieID, rating, timestamp) = line.split('\t')  
    yield rating, 1
```

# Writing the Reducer

USER ID|MOVIE ID|RATING|TIMESTAMP

196 242 3 881250949  
186 302 3 891717742  
196 377 1 878887116  
244 51 2 880606923  
166 346 1 886397596  
186 474 4 884182806  
186 265 2 881171488

Map

3,1  
3,1  
1,1  
2,1  
1,1  
4,1  
2,1

Shuffle  
& Sort

1 -> 1, 1  
2 -> 1, 1  
3 -> 1, 1  
4 -> 1

Reduce

1, 2  
2, 2  
3, 2  
4, 1

```
def reducer_count_ratings(self, key, values):  
    yield key, sum(values)
```

# Putting it all together


```
from mrjob.job import MRJob
from mrjob.step import MRStep

class RatingsBreakdown(MRJob):
    def steps(self):
        return [
            MRStep(mapper=self.mapper_get_ratings,
                  reducer=self.reducer_count_ratings)
        ]

    def mapper_get_ratings(self, _, line):
        (userID, movieID, rating, timestamp) = line.split('\t')
        yield rating, 1

    def reducer_count_ratings(self, key, values):
        yield key, sum(values)

if __name__ == '__main__':
    RatingsBreakdown.run()
```

A decorative L-shaped frame made of thick dark blue lines. One part of the frame runs vertically down the left side, and the other part runs horizontally across the bottom, meeting at a right angle in the bottom right corner.

# RUNNING MAPREDUCE WITH MRJOB

Run our MapReduce job in our Hadoop installation

# Installing what we need

## ■ PIP

- *Utility for installing Python packages*
- `su root`  
`yum install python-pip`

## ■ Nano

- `yum install nano`

## ■ MRJob

- `pip install mrjob`  
`exit` downgrade the PyYAML package to 5.4.1 --> `sudo pip install PyYAML==5.4.1`

## ■ Data files and the script

- `wget http://media.sundog-soft.com/hadoop/ml-100k/u.data`
- `wget http://media.sundog-soft.com/hadoop/RatingsBreakdown.py`

# Running with mrjob

- Run locally
  - *python RatingsBreakdown.py u.item*
- Run with Hadoop
  - *python MostPopularMovie.py -r hadoop --hadoop-streaming-jar /usr/hdp/current/hadoop-mapreduce-client/hadoop-streaming.jar u.data*





# YOUR CHALLENGE

Sort movies by popularity with Hadoop



# Challenge exercise

- Count up ratings given for each movie
  - *All you need is to change one thing in the mapper - we don't care about ratings now, we care about movie ID's!*
  - *Start with this and make sure you can do it.*
  - *You can use nano to just edit the existing RatingsBreakdown.py script*

# Stretch goal

- Sort the movies by their numbers of ratings
- Strategy:
  - *Map to (movieID, 1) key/value pairs*
  - *Reduce with output of (rating count, movieID)*
  - *Send this to a second reducer so we end up with things sorted by rating count!*
- Gotchas:
  - *How do we set up more than one MapReduce step?*
  - *How do we ensure the rating counts are sorted properly?*

# Multi-stage jobs

- You can chain map/reduce stages together like this:

```
def steps(self):  
    return [  
        MRStep mapper=self.mapper_get_ratings,  
            reducer=self.reducer_count_ratings),  
        MRStep(reducer=self.reducer_sorted_output)  
    ]
```

# Ensuring proper sorting

- By default, streaming treats all input and output as strings. So things get sorted as strings, not numerically.
- There are different formats you can specify. But for now let's just zero-pad our numbers so they'll sort properly.

- The second reducer will look like this:

```
def reducer_count_ratings(self, key, values):  
    yield str(sum(values)).zfill(5), key
```

# Iterating through the results

- Spoiler alert!

```
def reducer_sorted_output(self, count, movies):  
    for movie in movies:  
        yield movie, count
```

# CHECK YOUR RESULTS

Did it work?



# My solution

```
from mrjob.job import MRJob
from mrjob.step import MRStep

class RatingsBreakdown(MRJob):
    def steps(self):
        return [
            MRStep(mapper=self.mapper_get_ratings,
                    reducer=self.reducer_count_ratings),
            MRStep(reducer=self.reducer_sorted_output)
        ]

    def mapper_get_ratings(self, _, line):
        (userID, movieID, rating, timestamp) = line.split('\t')
        yield movieID, 1

    def reducer_count_ratings(self, key, values):
        yield str(sum(values)).zfill(5), key

    def reducer_sorted_output(self, count, movies):
        for movie in movies:
            yield movie, count

if __name__ == '__main__':
    RatingsBreakdown.run()
```