



APACHE DRILL

SQL for noSQL



What is Apache Drill?

- A SQL query engine for a variety of non-relational databases and data files
 - *Hive, MongoDB, HBase* Not Cassandra
 - *Even flat JSON or Parquet files on HDFS, S3, Azure, Google cloud, local file system*
- Based on Google's Dremel



Big Table --> HBase

GFS --> HDFS

Dremel --> Drill

It's real SQL

- Not SQL-Like
- And it has a ODBC / JDBC driver so other tools can connect to it just like any relational database



It's fast and pretty easy to set up.

- But remember, these are still non-relational databases under the hood!
- Allows SQL analysis of disparate data source without having to transform and load it first
 - *Internally data is represented as JSON and so has no fixed schema*

You can even do joins across different database technologies

- Or with flat JSON files that are just sitting around

Think of it as SQL for your entire ecosystem



Let's drill

- We'll import data into Hive and MongoDB
- Set up Drill on top of both
- And do some queries!

